

# Batch processing for smart plugs sensor data with Apache Spark on Google DataProc

Ovidiu Daniel Barba  
Laura Trivelloni  
Emanuele Vannacci




1.

# System Architecture

High Level Overview

# System Layers



# Data Processing Layer



- ◆ Basic RDD API
- ◆ Scala as the programming language
  - a. Spark written in Scala
  - b. Functional Programming Functions similar with RDD API


# High Level Interface



- ◆ Support for structured and semi-structured data using DataFrame API
- ◆ SQL-like query language
- ◆ Our own UDF (User Defined Functions)



ALLUXIO



***hadoop***



mongoDB®

# Data Storage

- ◆ HDFS stores and serves input, results and benchmark files
- ◆ Alluxio handles all data read, write and cache operations on HDFS theoretically improving performance up to 10x

# Data Ingestion



- ◆ Inject data from external sources to data storage layer
- ◆ Transfer data between components in storage layer
- ◆ Data filtering

# Data ingestion from source

**Navigate**

**Operate**

NiFi Flow  
Process Group


bae89b9d-0163-1000-d5b5-5b1fa3baee1f


DELETE





# Internal Data Transfer

 **Navigate**






 **Operate**



**NiFi Flow**  
Process Group






bae89b9d-0163-1000-d5b5-5b1fa3baee1f













DELETE

**ExtractFromHDFS**



3 0 0

Queued	0 (0 bytes)
In	0 (0 bytes) → 0
Read/Write	0 bytes / 0 bytes
Out	1 → 0 (0 bytes)








0 0 0 0 0 0 0

**From toRoute**

Queued	0 (0 bytes)
--------	-------------

5 min

**RouteOnQuery**  
RouteOnAttribute 1.6.0  
org.apache.nifi - nifi-standard-nar



0 0 0

In	0 (0 bytes)
Read/Write	0 bytes / 0 bytes
Out	0 (0 bytes)

Tasks/Time 0 / 00:00:00.000

5 min

**Name query1**

Queued	0 (0 bytes)
--------	-------------

5 min

**Name query2**

Queued	0 (0 bytes)
--------	-------------





5 min

**Name times**

Queued	0 (0 bytes)
--------	-------------

5 min

**PutMongoQ1**  
PutMongo 1.6.0  
org.apache.nifi - nifi-mongodb-nar








0 0 0

In	0 (0 bytes)
Read/Write	0 bytes / 0 bytes
Out	0 (0 bytes)

Tasks/Time 0 / 00:00:00.000

5 min

**PutMongoQ2**  
PutMongo 1.6.0  
org.apache.nifi - nifi-mongodb-nar



0 0 0

In	0 (0 bytes)
Read/Write	0 bytes / 0 bytes
Out	0 (0 bytes)

Tasks/Time 0 / 00:00:00.000

5 min

**PutMongoQ3**  
PutMongo 1.6.0  
org.apache.nifi - nifi-mongodb-nar







0 0 0

In	0 (0 bytes)
Read/Write	0 bytes / 0 bytes
Out	0 (0 bytes)

Tasks/Time 0 / 00:00:00.000

5 min

**PutMongoTimes**  
PutMongo 1.6.0  
org.apache.nifi - nifi-mongodb-nar



0 0 0

In	0 (0 bytes)
Read/Write	0 bytes / 0 bytes
Out	0 (0 bytes)

Tasks/Time 0 / 00:00:00.000

5 min

# Data provenance lineage

## Data ingestion from local to HDFS < 30 sec

06/05/2018 20:09:47.929 UTC	CREATE	d782b814-b23a-4ce0-bc63-e5ee06c08ea	0 bytes	ListFile
...				
06/05/2018 20:10:05.687 UTC	DROP	d6cf3e15-a6a8-42e4-8a28-234d42b58122	56.84 MB	PutHDFS
06/05/2018 20:10:08.460 UTC	DROP	bc9abcf9-1dc7-4495-ada2-bbdee0186ac7	12.28 MB	PutHDFS
06/05/2018 20:10:12.517 UTC	DROP	6d650841-fc56-4687-85ca-a40d13d2550c	56.84 MB	PutParquet

## Results collection from HDFS to MongoDB < 10 sec

06/06/2018 10:05:21.342 UTC	CREATE	bd55a8a3-a42e-470d-b40a-d35f...	0 bytes	ListFile	ListFile
...					
06/06/2018 10:05:30.146 UTC	DROP	87622296-7d1e-49f2-b789-3a6a...	2.41 KB	PutMongoQ3	PutMongo

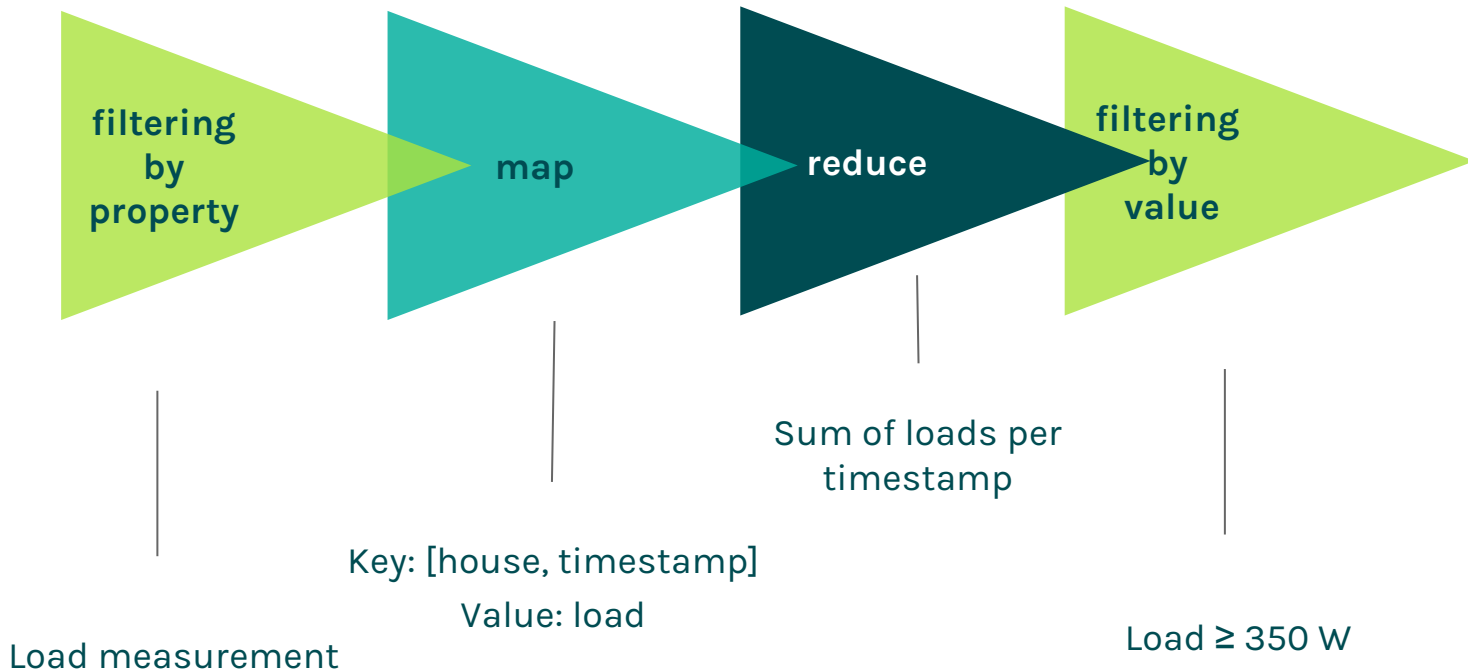


2.

# Queries

Detailed queries description


# Query 1



## Query 1 Spark SQL

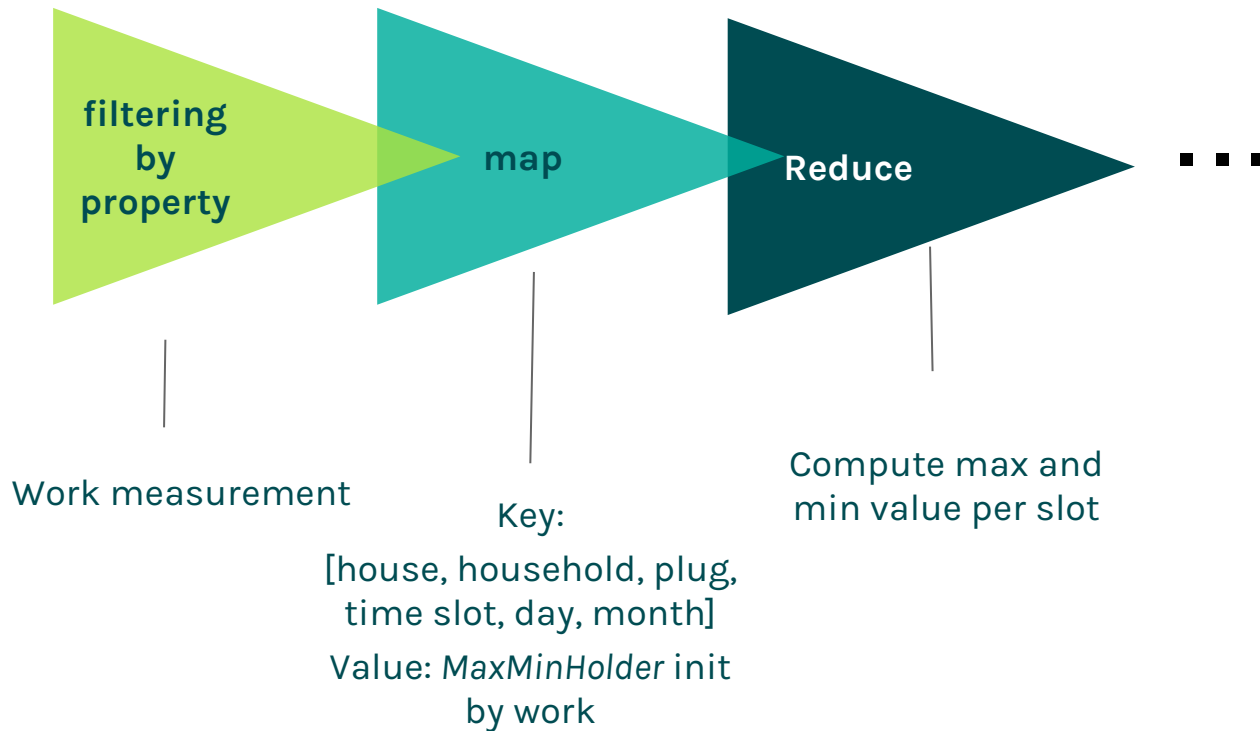
```
val res = df
  .where("property = 1")
  .groupBy("house_id", "timestamp")
  .agg(sum("value").as("sum"))
  .select("house_id")
  .where("sum >= 350")
  .distinct()
  .sort($"house_id")
  .collect()
```

# Query 1 performances

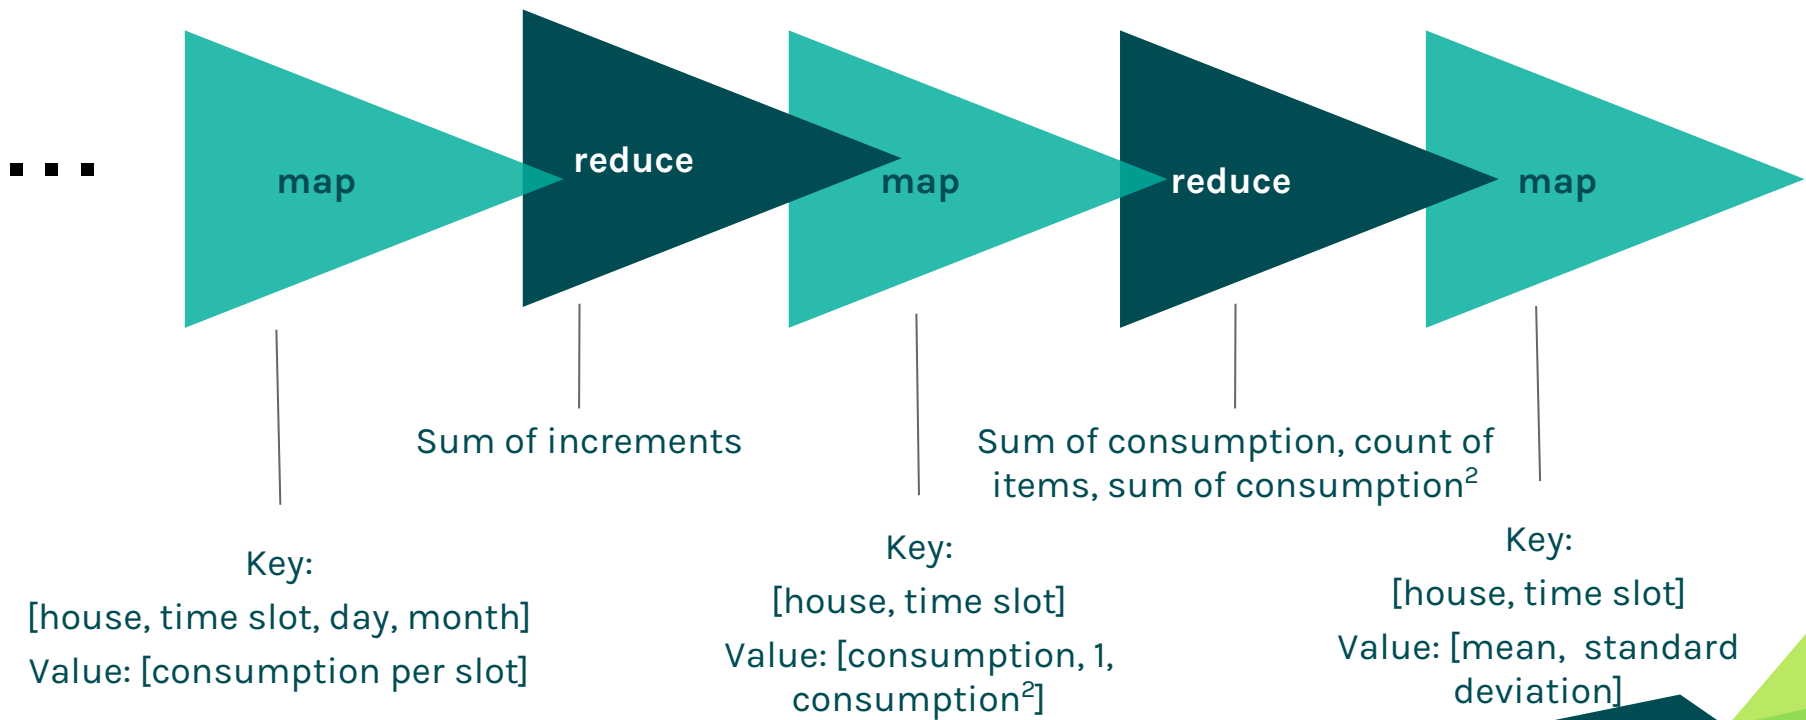


Spark Core and Spark SQL implementations performances.

## Query 2



## Query 2





## Query 2 Spark SQL

```
val data = df
  .where("property == 0")
  .withColumn("timestamp", to_utc_timestamp(from_unixtime($"timestamp"), "Etc/GMT+2"))
  .withColumn("value", $"value".cast(DataTypes.createDecimalType(20, 5)))


  .groupBy($"house_id", $"household_id", $"plug_id", window($"timestamp", "6 hours"))
  .agg(
    when(last("value") >= first("value"), last("value") - first("value"))
      .otherwise(last("value"))
      .alias("plug_consumption")
  )

  .groupBy("house_id", "window")
  .agg(sum($"plug_consumption").as("home_consumption"))
  .withColumn("window", struct(date_format($"window.start", "HH:mm"), date_format($"window.end", "HH:mm")))

  .groupBy($"house_id", $"window")
  .agg(avg("home_consumption").as("avg"), stddev("home_consumption").as("stddev"))

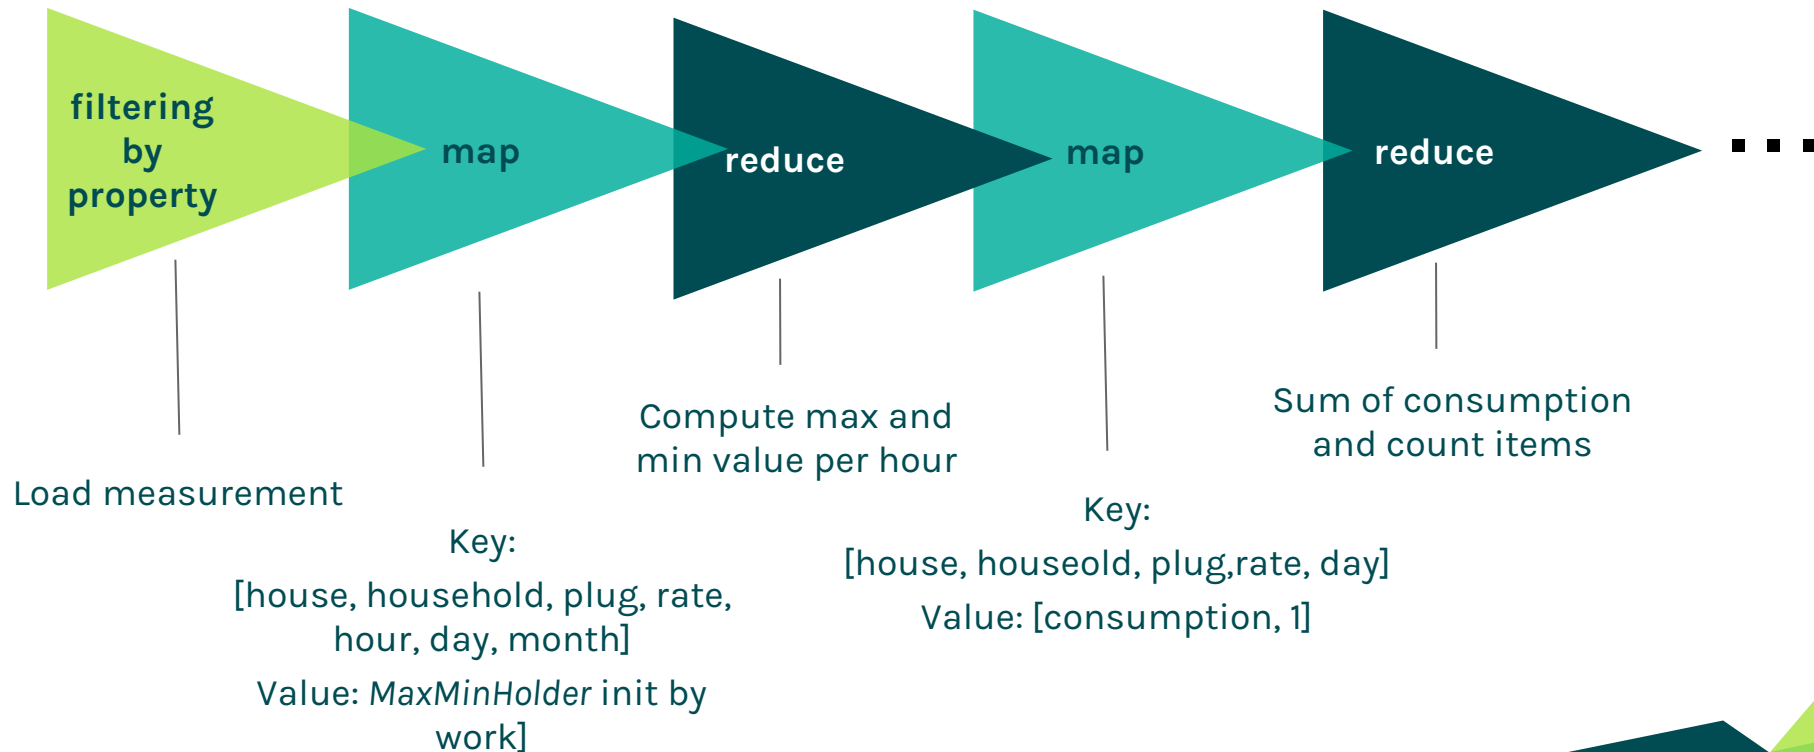
  .orderBy("house_id", "window")
  .select("*")
  .collect()
```

# Query 2 performances

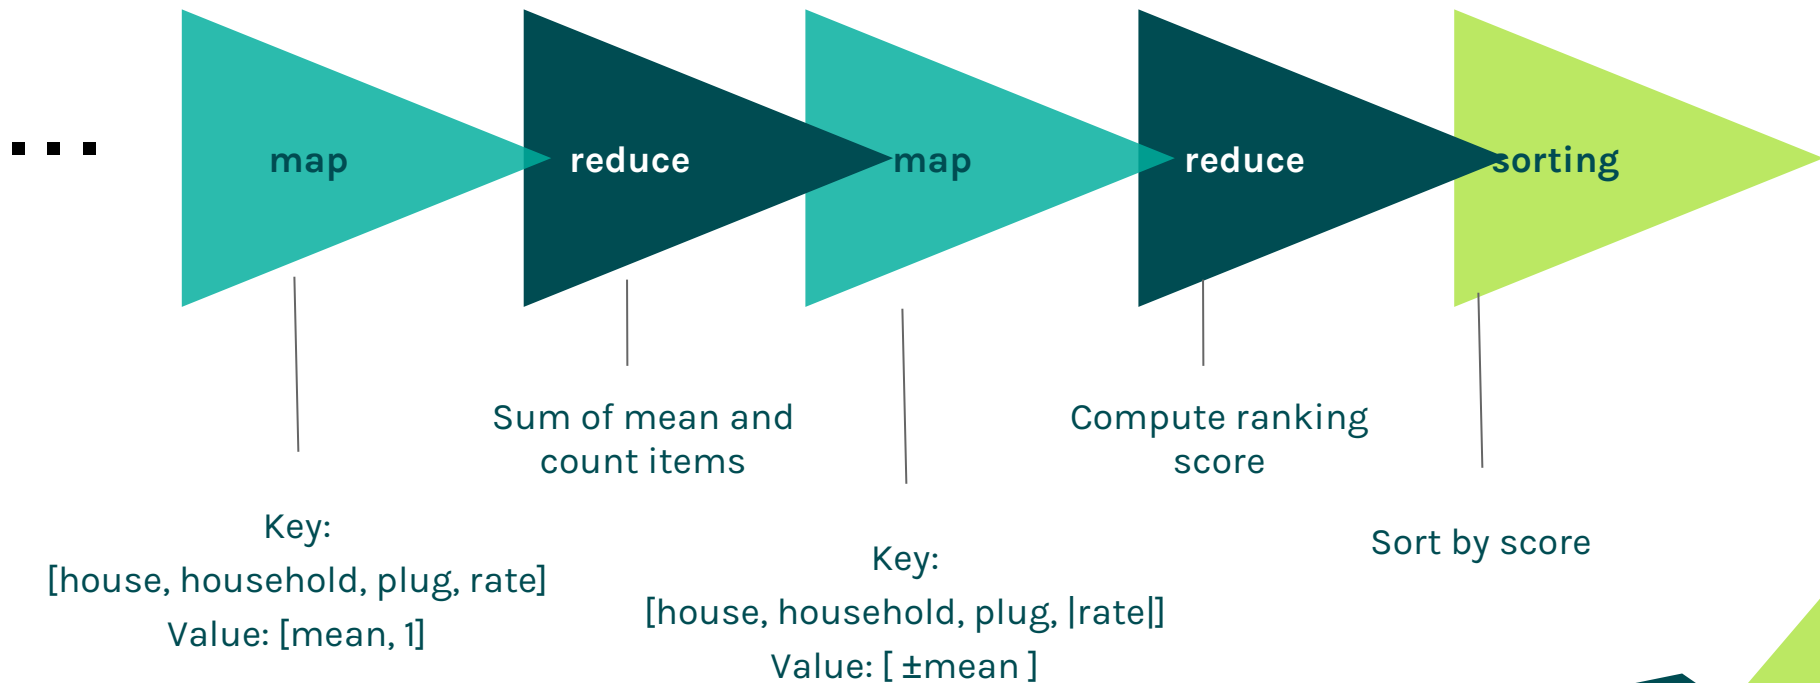


Spark Core and Spark SQL implementations performances.

## Query 3



## Query 3



## Query 3 Spark SQL

```
val res = df
  .where("property = 0")
  .withColumn("value", $"value".cast(DataTypes.createDecimalType(20, 5)))
  .withColumn("slot", udfDataFunction.getPeriodRateUDF('timestamp))
  .withColumn("day", udfDataFunction.getDayOfMonthUDF('timestamp))

  .groupBy($"house_id", $"household_id", $"plug_id",
    udfDataFunction.getHourOfDayUDF('timestamp), $"day", $"slot")
  .agg((max("value") - min("value")).as("plug_consumption_hour"))


  .groupBy($"house_id", $"household_id", $"plug_id", $"day", $"slot")
  .agg(avg("plug_consumption_hour").as("plug_consumption_day"))

  .groupBy($"house_id", $"household_id", $"plug_id", $"slot")
  .agg(avg("plug_consumption_day").as("avg"))
  .withColumn("avg", udfDataFunction.invertSignUDF('avg, 'slot))

  .groupBy($"house_id", $"household_id", $"plug_id", abs($"slot").as("month"))
  .agg(sum("avg").as("score"))

  .orderBy(desc("score"))
  .collect()
```


# Query 3 performances



Spark Core and Spark SQL implementations performances.

# Query 3 improvement

## Implementation improvement





3.

# Results

Queries output on given dataset



## Query 1 results

House ID	0	1	2	3	4	5	6	7	8	9
----------	---	---	---	---	---	---	---	---	---	---

## Query 2 (partial) results

House ID	0	0	0	0	1	1	1	1	2	2	2	2
Time slot	00:00-5:59	6:00-11:59	12:00-17:59	18:00-23:59	00:00-5:59	6:00-11:59	12:00-17:59	18:00-23:59	00:00-5:59	6:00-11:59	12:00-17:59	18:00-23:59
Mean	0.14	0.41	0.19	0.25	0.37	0.69	0.79	0.68	0.30	0.30	0.31	0.33
Standard deviation	0.11	1.04	0.14	0.14	0.36	0.48	0.49	0.40	0.31	0.32	0.30	0.33

...

## Query 3 (partial) results

Ranking	Plug	Month	Score
1°	(8,0,1)	9	0.115
2°	(8,0,0)	9	0.080
3°	(0,0,2)	9	0.077
4°	(1,0,1)	9	0.042
5°	(5,0,1)	9	0.011
6°	(7,0,1)	9	0.008

...




4.


# Deployment

Local and Cloud deployment

# Local Deploy



- ◆ Every system component is built as a Docker Image
- ◆ Made from scratch (Spark, NiFi, Alluxio) or from existing images (HDFS, MongoDB)
- ◆ Some run in pseudo-distributed mode (HDFS, Spark and Alluxio) with 1 Master and N Workers
- ◆ All components run on same network



# Docker Images



ovidanb/alluxio  
public

0  
STARS

65  
PULLS

>  
DETAILS



ovidanb/spark  
public

0  
STARS

1  
PULLS

>  
DETAILS



zanna94/nifi  
public

0  
STARS

37  
PULLS

>  
DETAILS

# Cloud Deploy

Google Cloud Platform was used as the deploy environment.

In particular the **DataProc** and **Kubernetes Engine** services.



Google Cloud Platform

# DataProc

- ◆ Fully-managed cloud service for running Spark and Hadoop Clusters
- ◆ Ready with Hadoop Ecosystem frameworks (HDFS, Hive, Pig and YARN)
- ◆ 3 or 5 node cluster configurations (1 Master only)



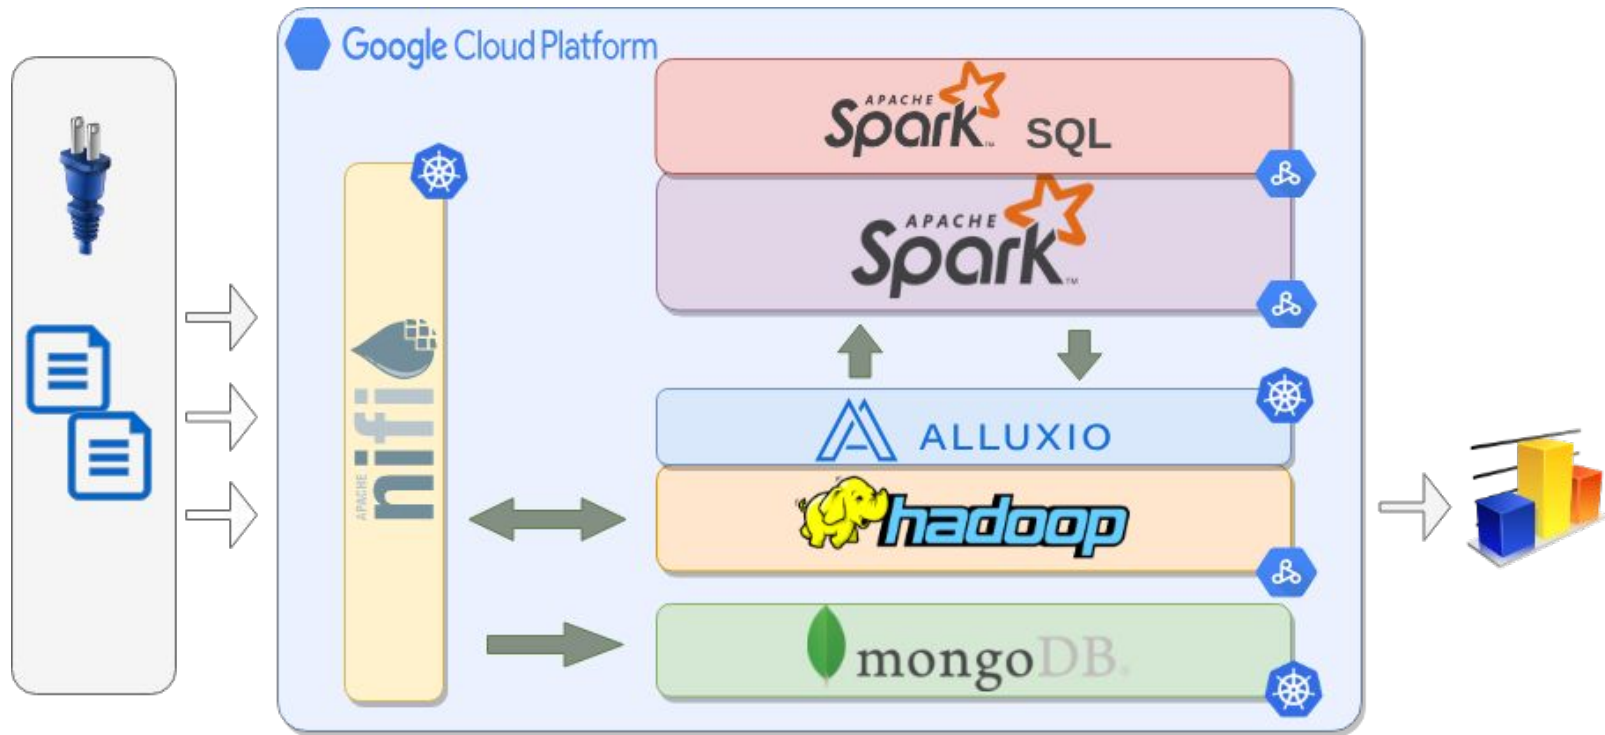


# Kubernetes



- ◆ Used to deploy remaining System components (MongoDB, NiFi and Alluxio) as Docker containers
- ◆ 3-node Cluster (1 Master and 2 Slaves)
- ◆ On the same VPC network as DataProc cluster

# Cloud Deployment



# Spark CPU Utilization

VM instance details

EDIT

RESET

CLONE

STOP

DELETE

Details Monitoring

Reset zoom

**1 hour**

6 hours

12 hours

1 day

2 days

4 days

7 days


14 days

30 days

CPU

% CPU


Jun 4, 2018 5:27 PM



Network Bytes

Bytes/sec

Jun 4, 2018 5:27 PM



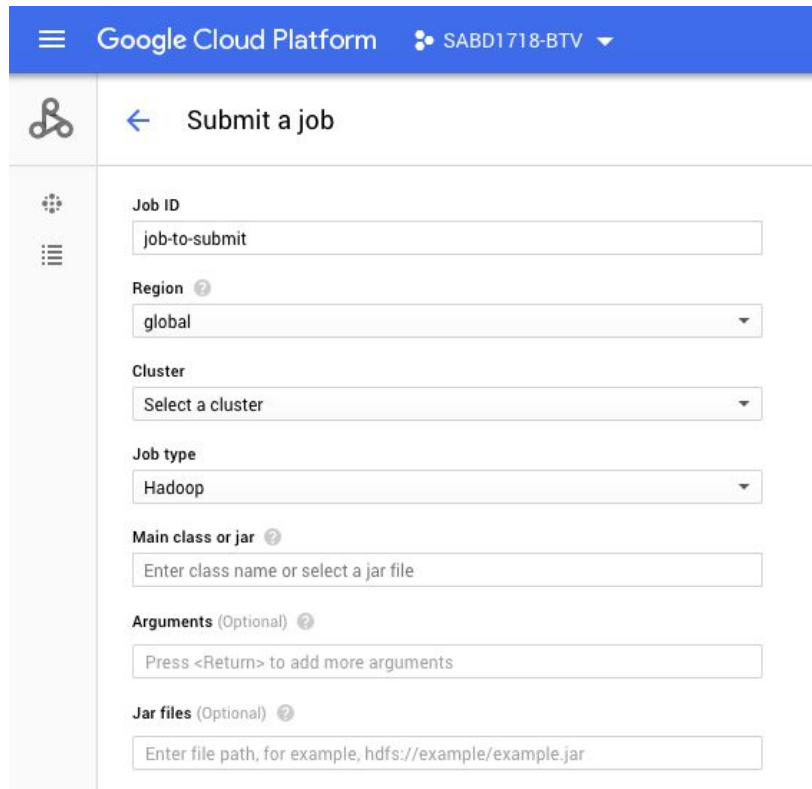
Network Packets

Packets/sec

# DataProc Submit Job

## Job Type:

- ◆ Hadoop
- ◆ Spark
- ◆ PySpark
- ◆ Hive
- ◆ SparkSQL
- ◆ Pig



The screenshot shows the 'Submit a job' page in the Google Cloud Platform DataProc console. The page has a blue header with the Google Cloud Platform logo and the account 'SABD1718-BTV'. The main content area is white with a left sidebar containing navigation icons. The form fields are as follows:

- Job ID:** A text input field containing 'job-to-submit'.
- Region:** A dropdown menu with 'global' selected.
- Cluster:** A dropdown menu with 'Select a cluster' selected.
- Job type:** A dropdown menu with 'Hadoop' selected.
- Main class or jar:** A text input field with the placeholder 'Enter class name or select a jar file'.
- Arguments (Optional):** A text input field with the placeholder 'Press <Return> to add more arguments'.
- Jar files (Optional):** A text input field with the placeholder 'Enter file path, for example, hdfs://example/example.jar'.

# DataProc Jobs - 1

Jobs

[+ SUBMIT JOB](#)

[REFRESH](#)

[STOP](#)

[DELETE](#)

[REGIONS](#)

<input type="checkbox"/>		job-smart-plug-run-hdfs-3	europe-west4	Spark	cluster-spark	Jun 1, 2018, 1:29:34 PM	20 min 37 sec	Succeeded
<input type="checkbox"/>		job-smart-plug-run-hdfs-2	europe-west4	Spark	cluster-spark	Jun 1, 2018, 1:10:19 PM	10 min 37 sec	Succeeded
<input type="checkbox"/>		job-smart-plug-run-hdfs-1	europe-west4	Spark	cluster-spark	Jun 1, 2018, 12:42:58 PM	10 min 34 sec	Succeeded
<input type="checkbox"/>		job-smart-plug-run-4	europe-west4	Spark	cluster-spark	May 31, 2018, 11:11:15 PM	1 hr 45 min	Succeeded
<input type="checkbox"/>		job-smart-plug-run-3	europe-west4	Spark	cluster-spark	May 31, 2018, 8:55:52 PM	50 min 0 sec	Succeeded
<input type="checkbox"/>		job-smart-plug-run-2	europe-west4	Spark	cluster-spark	May 31, 2018, 8:54:00 PM	2 min 37 sec	Canceled
<input type="checkbox"/>		job-smart-plug-run-1	europe-west4	Spark	cluster-spark	May 31, 2018, 6:56:41 PM	53 min 51 sec	Succeeded
<input type="checkbox"/>		job-smart-plug-8	europe-west4	Spark	cluster-spark	May 31, 2018, 5:58:38 PM	39 min 54 sec	Failed
<input type="checkbox"/>		job-smart-plug-7	europe-west4	Spark	cluster-spark	May 31, 2018, 5:41:11 PM	1 min 6 sec	Failed
<input type="checkbox"/>		job-smart-plug-6	europe-west4	Spark	cluster-spark	May 31, 2018, 5:12:06 PM	1 min 1 sec	Failed
<input type="checkbox"/>		job-smart-plug-5	europe-west4	Spark	cluster-spark	May 31, 2018, 5:04:32 PM	59 sec	Failed
<input type="checkbox"/>		job-smart-plug-4	europe-west4	Spark	cluster-spark	May 31, 2018, 4:55:09 PM	57 sec	Failed
<input type="checkbox"/>		job-smart-plug-3	europe-west4	Spark	cluster-spark	May 31, 2018, 4:25:15 PM	1 min 17 sec	Failed
<input type="checkbox"/>		job-smart-plug-2	europe-west4	Spark	cluster-spark	May 31, 2018, 4:18:28 PM	24 sec	Failed
<input type="checkbox"/>		job-smart-plug-1	europe-west4	Spark	cluster-spark	May 31, 2018, 4:12:41 PM	5 sec	Failed
<input type="checkbox"/>		job-smart-plug	europe-west4	Spark	cluster-spark	May 31, 2018, 4:12:06 PM	4 sec	Failed
<input type="checkbox"/>		demo-2	europe-west4	Spark	cluster-spark	May 29, 2018, 10:10:19 PM	30 sec	Succeeded
<input type="checkbox"/>		job-3	europe-west4	Spark	cluster-spark	May 29, 2018, 10:04:09 PM	15 sec	Succeeded
<input type="checkbox"/>		job-2	europe-west4	Spark	cluster-spark	May 29, 2018, 10:03:09 PM	12 sec	Failed
<input type="checkbox"/>		job-6aa546fc	europe-west4	Spark	cluster-spark	May 29, 2018, 9:42:38 PM	15 sec	Failed
<input type="checkbox"/>		demo-1	europe-west4	Spark	plug-cluster	May 28, 2018, 8:50:25 PM	53 sec	Succeeded

# DataProc Jobs - 2

Jobs

[+ SUBMIT JOB](#)

[REFRESH](#)

[STOP](#)

[DELETE](#)

[REGIONS](#)

Search jobs, press Ente

<input type="checkbox"/>	Job ID	Region	Type	Cluster	Start time	Elapsed time	Status
<input type="checkbox"/>	✔ job-smart-plug-5w-object-alluxio-11	europe-west4	Spark	cluster-spark-large	Jun 4, 2018, 8:08:31 PM	22 min 30 sec	Succeeded
<input type="checkbox"/>	✔ job-smart-plug-5w-object-alluxio-10	europe-west4	Spark	cluster-spark-large	Jun 4, 2018, 7:53:41 PM	13 min 6 sec	Succeeded
<input type="checkbox"/>	✔ job-smart-plug-5w-object-9	europe-west4	Spark	cluster-spark-large	Jun 4, 2018, 7:36:31 PM	14 min 15 sec	Succeeded
<input type="checkbox"/>	✔ job-smart-plug-5w-object-8	europe-west4	Spark	cluster-spark-large	Jun 4, 2018, 6:52:12 PM	12 min 45 sec	Succeeded
<input type="checkbox"/>	✔ job-smart-plug-app-object-7	europe-west4	Spark	cluster-spark	Jun 4, 2018, 5:25:58 PM	30 min 19 sec	Succeeded
<input type="checkbox"/>	✔ job-smart-plug-app-object-6	europe-west4	Spark	cluster-spark	Jun 4, 2018, 4:38:47 PM	39 min 31 sec	Succeeded
<input type="checkbox"/>	❗ job-smart-plug-app-object-5	europe-west4	Spark	cluster-spark	Jun 4, 2018, 4:37:32 PM	31 sec	Failed
<input type="checkbox"/>	❗ job-smart-plug-app-object-4	europe-west4	Spark	cluster-spark	Jun 4, 2018, 4:35:03 PM	25 sec	Failed
<input type="checkbox"/>	✔ job-smart-plug-app-object-3	europe-west4	Spark	cluster-spark	Jun 4, 2018, 12:35:21 PM	32 min 26 sec	Succeeded
<input type="checkbox"/>	✔ job-smart-plug-app-random-2	europe-west4	Spark	cluster-spark	Jun 4, 2018, 12:33:48 PM	10 min 20 sec	Succeeded
<input type="checkbox"/>	✔ job-smart-plug-app-normal-1	europe-west4	Spark	cluster-spark	Jun 4, 2018, 12:33:06 PM	13 min 40 sec	Succeeded
<input type="checkbox"/>	✔ job-smart-plug-app-random-1	europe-west4	Spark	cluster-spark	Jun 4, 2018, 11:00:18 AM	7 min 29 sec	Succeeded
<input type="checkbox"/>	✔ job-smart-plug-app-object-2	europe-west4	Spark	cluster-spark	Jun 4, 2018, 10:52:08 AM	9 min 39 sec	Succeeded
<input type="checkbox"/>	✔ job-smart-plug-app-object-1	europe-west4	Spark	cluster-spark	Jun 4, 2018, 10:41:28 AM	7 min 24 sec	Succeeded
<input type="checkbox"/>	✔ job-smart-plug-run-alluxio-3	europe-west4	Spark	cluster-spark	Jun 1, 2018, 2:55:28 PM	20 min 39 sec	Succeeded
<input type="checkbox"/>	✔ job-smart-plug-run-alluxio-2	europe-west4	Spark	cluster-spark	Jun 1, 2018, 2:41:33 PM	9 min 40 sec	Succeeded
<input type="checkbox"/>	❗ job-smart-plug-run-alluxio-1	europe-west4	Spark	cluster-spark	Jun 1, 2018, 2:40:45 PM	12 sec	Failed
<input type="checkbox"/>	✔ job-smart-plug-run-hdfs-3	europe-west4	Spark	cluster-spark	Jun 1, 2018, 1:29:34 PM	20 min 37 sec	Succeeded
<input type="checkbox"/>	✔ job-smart-plug-run-hdfs-2	europe-west4	Spark	cluster-spark	Jun 1, 2018, 1:10:19 PM	10 min 37 sec	Succeeded
<input type="checkbox"/>	✔ job-smart-plug-run-hdfs-1	europe-west4	Spark	cluster-spark	Jun 1, 2018, 12:42:58 PM	10 min 34 sec	Succeeded

# Job Example - 1

## Job details

[REFRESH](#)[CLONE](#)

✓ job-smart-plug-5w-object-8

Start time: Jun 4, 2018, 6:52:12 PM Elapsed time: 12 min 45 sec Status: Succeeded

Output [Configuration](#)

[Edit](#)

Region	europe-west4
Cluster	cluster-spark-large
Job type	Spark
Main class or jar	BenchmarkMain
Jar files	gs://app_tvb/app_object.jar

Properties

Arguments

hdfs://35.204.61.244:8020/results-object-8
hdfs://35.204.61.244:8020/alluxio/data.csv
hdfs://35.204.61.244:8020/alluxio/data.parquet
hdfs://35.204.61.244:8020/alluxio/data.avro
cluster
cache
10

Labels

[+ Add label](#)

Equivalent [REST](#)

# Job Example - 2

## Job details

[REFRESH](#)[CLONE](#)

✓ job-smart-plug-5w-object-alluxio-11

Start time: Jun 4, 2018, 8:08:31 PM Elapsed time: 22 min 30 sec Status: Succeeded

Output [Configuration](#)

[Edit](#)

Region	europe-west4
Cluster	cluster-spark-large
Job type	Spark
Main class or jar	BenchmarkMain
Jar files	gs://app_tvb/app_object.jar

### Properties

#### Arguments

alluxio://35.204.176.216:30998/results-object-11  
alluxio://35.204.176.216:30998/data.csv  
alluxio://35.204.176.216:30998/data.parquet  
alluxio://35.204.176.216:30998/data.avro  
cluster  
cache  
10

#### Labels

[+ Add label](#)

Equivalent [REST](#)





5.

# Queries Performance

Execution time with different formats on  
different architectures


# Local Deployment Performance

## Specs :

- ◆ Macbook Air 2014
- ◆ 4 GB RAM
- ◆ 1,4 GHz Intel Core i5

## Docker (same network):


- 4 HDFS nodes
- 3 Spark nodes
- 2 Alluxio nodes



# Cloud deployment using HDFS

## DataProc :


- ◆ 5 worker node cluster
- ◆ 1 master node
- ◆ node: n1-standard-2 (2 vCPUs with 7.5 GB RAM)



# Cloud deployment using HDFS and cache

## DataProc :

- ◆ 5 worker node cluster
- ◆ 1 master node
- ◆ node: n1-standard-2 (2 vCPUs with 7.5 GB RAM)



# Cloud deployment using Alluxio


## DataProc :

- ◆ 5 worker node cluster & 1 master
- ◆ node: n1-standard-2 (2 vCPUs with 7.5 GB RAM)

## Kubernetes:

- 3 node n1-standard-2 cluster (each with Alluxio agent)

Both clusters on same VPC network



# Best Performance

- ◆ Avro as dataset format
- ◆ Alluxio (on top of HDFS) as distributed file system
- ◆ No Spark caching



# Best Performance

	Core	SQL
Query 1	2.37 s	2,01 s
Query 2	1.39 s	4.12 s
Query 3	2.84 s	5.31 s



ALLUXIO

# Thank you for listening!

